

# Hierarchical Strategic Coordination for Infrastructure-Aware Wildfire Suppression under Realistic Physics

Anonymous submission

## Abstract

Coordinating multi-agent teams for wildfire suppression under realistic fire physics is a severe spatial-temporal challenge. Standard decentralized multi-agent reinforcement learning (MARL) suffers from *coordination collapse*: agents with limited egocentric observations redundantly suppress nearby cells while failing to defend distant, high-value infrastructure. We propose a decoupled hierarchical architecture that separates *strategic macro-dispatch*—assigning agents to spatial sectors based on threat and asset priority—from *tactical micro-execution* via a deterministic target-seeking compliance policy. To stabilize strategic learning, we warm-start the commander via behavior cloning (BC) from a domain-expert heuristic and fine-tune with KL-regularized policy gradients and infrastructure-weighted economic rewards. We evaluate on two real-world GIS topographies using validated Cell2Fire wildfire physics: Saudi Arabian petroleum infrastructure and California Wildland-Urban Interface (WUI) zones. Over  $5 \text{ seeds} \times 15 \text{ episodes}$ , our learned hierarchy statistically matches a hand-crafted expert heuristic on Weighted Economic Loss ( $p = 0.115$  Saudi;  $p = 0.952$  California) while dramatically outperforming flat MARL ( $p = 3.0 \times 10^{-5}$ , Cohen’s  $d = 13.28$ ). Compliance diagnostics confirm near-perfect target adherence (compliance rate 1.000, dispatch latency 1.67 steps), providing interpretable evidence that strategic intent translates into coordinated physical behavior. Code, checkpoints, and reproducibility manifests are available in the supplementary material.

## 1 Introduction

Wildfire suppression demands coordinated allocation of limited suppression resources across a dynamically evolving hazard landscape. When critical infrastructure—petroleum refineries, pipeline junctions, residential structures—lies in the path of an advancing fire front, the coordination challenge intensifies: agents must *strategically prioritize* which assets to defend rather than simply minimizing total burned area. This infrastructure-defense framing distinguishes wildfire suppression from generic fire-containment problems and motivates the present work.

This is an anonymized submission for review purposes only. Distribution, citation, or public sharing of this manuscript is strictly prohibited. Copyright and publication details will appear in the final version if accepted.

**Problem: Flat MARL Fails.** Standard decentralized MARL algorithms such as MAPPO (Yu et al. 2022) and QMIX (Rashid et al. 2018) equip agents with local egocentric observations and learn decentralized policies via centralized training. Under partial observability and dynamic fire spread, these methods exhibit *coordination collapse*: agents cluster near locally visible fire, redundantly treating the same cells while distant high-value assets burn undefended. In our experiments, flat MARL agents achieve infrastructure survival rates indistinguishable from a no-operation baseline (Section 6).

**Insight: Strategic Intent Requires Tactical Compliance.** Hierarchical reinforcement learning (HRL) can, in principle, resolve this failure by introducing a strategic commander that dispatches agents to priority sectors. However, a critical systems-level issue arises: if low-level agents do not reliably *comply* with strategic directives, the commander’s decisions become decoupled from physical outcomes. We term this the *compliance gap*—the discrepancy between strategic intent and realized agent behavior. Prior HRL work rarely measures or reports this gap.

**Solution: Decoupled Hierarchical Coordination.** We propose a two-timescale architecture: (1) a learned **strategic commander** that assigns each agent to one of  $M = 16$  spatial sectors every  $T_{\text{macro}} = 10$  steps, and (2) a deterministic **target-seeking compliance policy** that routes agents to their assigned sector centers and executes local suppression. By decoupling tactical execution from strategic learning, we guarantee that the commander’s decisions directly control spatial allocation, eliminating the compliance gap.

**Result: Expert-Level Infrastructure Defense.** On two real-world GIS topographies—Saudi Arabian petroleum grids and California WUI zones—evaluated under validated Cell2Fire wildfire physics, our learned hierarchy statistically matches a domain-expert heuristic baseline on infrastructure-weighted economic metrics while dramatically outperforming flat decentralized MARL (Table 1).

**Contributions.** Our contributions are:

1. A **decoupled hierarchical MARL framework** for wildfire suppression built on validated Cell2Fire physics, providing infrastructure-aware benchmarks on real Saudi and California GIS topographies.

2. A **stable training pipeline** combining BC warm-start, KL-regularized policy gradients, and infrastructure-weighted economic reward shaping.
3. A **compliance diagnostics framework** measuring target adherence, dispatch latency, and sector drift—providing interpretable evidence that strategic intent translates into physical coordination.
4. **Rigorous statistical evaluation** across 5 seeds with bootstrap 95% CIs, Welch’s  $t$ -tests, Cohen’s  $d$  effect sizes, and component ablations.

## 2 Related Work

**Wildfire Simulation and GIS-Based Planning.** Wildfire spread modeling has progressed from simple cellular automata to physics-based simulators. Cell2Fire (Pais et al. 2021) implements the Canadian Forest Fire Behavior Prediction (FBP) system with heterogeneous fuel types, wind effects, and topographic slope corrections. Operational research approaches optimize firebreak placement using mixed-integer programming (Hu and Ntamo 2009) but assume global state observability and static fire behavior, neither of which holds during active suppression. Our work uses Cell2Fire as a validated physics backend while framing suppression as a sequential decision problem under partial observability.

**Reinforcement Learning for Wildfire and Disaster Response.** Recent RL approaches to wildfire include UAV-based detection (Ghali, Akhloufi, and Wael Saber 2022; Partheepan, Sanati, and Hassan 2023), landscape-scale fuel-treatment planning (Altamimi et al. 2022), and grid-based suppression agents on Cell2Fire (Shen and Curtis 2022). However, most use simplified fire dynamics without validated physics, single-agent formulations, or toy grid sizes that lack realistic infrastructure. Disaster-response RL more broadly has explored resource allocation (Nguyen, Nguyen, and Nahavandi 2020), but rarely addresses the spatial coordination challenge of multi-agent infrastructure defense under physically grounded fire spread.

**Cooperative MARL and Communication.** Centralized Training with Decentralized Execution (CTDE) is the dominant paradigm for cooperative MARL. QMIX (Rashid et al. 2018) factorizes joint value functions under monotonicity constraints; MAPPO (Yu et al. 2022) applies PPO (Schulman et al. 2017) with centralized critics. CommNet (Sukhbaatar, Szlam, and Fergus 2016) and TarMAC (Das et al. 2019) add learned communication channels. These methods assume agents can learn implicit coordination from reward signals alone. Under partial observability with dynamic wildfire spread, we show empirically that this assumption fails—flat MARL agents exhibit coordination collapse identical to no-operation baselines.

**Hierarchical Reinforcement Learning.** The options framework (Sutton, Precup, and Singh 1999) and feudal networks (Vezhnevets et al. 2017) introduce temporal abstraction via subgoal hierarchies. HAM (Parr and Russell 1997) and MAXQ (Dietterich 1998) decompose MDPs into sub-task hierarchies. In multi-agent settings, RODE (Wang et al.

2021) uses role decomposition, while MAVEN (Mahajan et al. 2019) conditions on shared latent variables. Most prior HRL work uses *learned* low-level policies that are jointly optimized with the high-level controller. A key distinction of our approach is that we *deliberately decouple* the tactical layer into a deterministic target-seeking policy, which guarantees that strategic decisions control spatial allocation and enables direct compliance measurement.

**Imitation Learning and RL Fine-Tuning.** Combining imitation with RL is well-established: DAGger (Ross, Gordon, and Bagnell 2011) iteratively corrects distributional shift; behavioral cloning followed by RL fine-tuning has been applied to game playing (Silver et al. 2016) and robotics (Rajeswaran et al. 2018). KL-regularized policy optimization prevents catastrophic forgetting of imitation priors (Wu, Tucker, and Nachum 2019). We apply this paradigm to hierarchical wildfire coordination, using a domain-expert value-weighted heuristic as the BC teacher for the strategic commander.

## 3 Problem Formulation

We model wildfire suppression as a Decentralized Partially Observable Markov Decision Process (Dec-POMDP) extended with hierarchical temporal abstraction.

**State.** The global state  $s = \langle S_{\text{fire}}, S_{\text{asset}}, S_{\text{agent}} \rangle$  consists of:

- $S_{\text{fire}} \in \{0, 1\}^{H \times W}$ : binary active-fire grid evolved by Cell2Fire physics.
- $S_{\text{asset}} \in \mathbb{R}_{\geq 0}^{H \times W}$ : infrastructure criticality map encoding economic asset values.
- $S_{\text{agent}} = \{(y_i, x_i)\}_{i=1}^N$ : grid positions of  $N$  suppression agents.

**Observations.** Each agent  $i$  observes an egocentric  $C \times C$  crop ( $C = 9$ ) containing 8 channels: active fire, treated cells, fuel type, criticality, agent occupancy, normalized coordinates, fire fraction, and time fraction. Agents do *not* observe the full grid.

**Actions.** Micro-actions  $a_i^t \in \{\text{stay, up, down, left, right, treat}\}$  are executed at every environment step. Macro-actions  $a_i^{\text{macro}} \in \{0, \dots, M-1\}$  assign agents to one of  $M = 16$  spatial sectors (a  $4 \times 4$  partition of the grid), selected every  $T_{\text{macro}} = 10$  steps.

**Reward.** The macro-level reward for the strategic commander is:

$$R^{\text{macro}} = -\Delta \text{WEL} + \alpha \cdot \Delta \text{ISR} \quad (1)$$

where  $\text{WEL} = \sum_j v_j \cdot 1[\text{asset}_j \text{ burned}]$  is the Weighted Economic Loss,  $\text{ISR} = |\{\text{surviving assets}\}|/|\{\text{total assets}\}|$  is the Infrastructure Survival Rate, and  $\alpha = 20$  upweights survival. This reward directly incentivizes infrastructure defense over generic burned-area minimization.




Figure 1: **Decoupled hierarchical architecture.** The strategic commander observes coarse  $4 \times 4$  grids of fire threat, asset criticality, and agent positions, producing sector assignments every  $T_{\text{macro}} = 10$  steps. Tactical agents execute deterministic target-seeking to reach assigned sector centers, then suppress locally. Arrows indicate the decoupled information flow.

## 4 Methodology

### 4.1 Strategic Commander

The commander  $\pi_\theta$  operates on coarse state representations obtained by average-pooling the  $H \times W$  grids into  $4 \times 4$  summaries:

$$\bar{S}_{\text{fire}} = \text{AvgPool}_{4 \times 4}(S_{\text{fire}}) \in \mathbb{R}^{4 \times 4} \quad (2)$$

$$\bar{S}_{\text{asset}} = \text{AvgPool}_{4 \times 4}(S_{\text{asset}}) \in \mathbb{R}^{4 \times 4} \quad (3)$$

$$\bar{S}_{\text{agent}} = \text{OneHot}_{4 \times 4}(S_{\text{agent}}) \in \{0, 1\}^{4 \times 4} \quad (4)$$

These three channels are concatenated and passed through a shared MLP encoder, producing per-agent logits over  $M = 16$  sectors:

$$\pi_\theta(a_i^{\text{macro}} | s) = \text{Softmax}(f_\theta(\bar{S}_{\text{fire}}, \bar{S}_{\text{asset}}, \bar{S}_{\text{agent}})_i) \quad (5)$$

### 4.2 Tactical Target-Seeking Compliance

Given a sector assignment  $a_i^{\text{macro}} = k$ , the tactical policy routes agent  $i$  toward the center  $(y_k^*, x_k^*)$  of sector  $k$  using deterministic Chebyshev-distance minimization. If the agent is within 1 cell of the center and an adjacent cell is actively burning, the agent executes the `treat` action. This design guarantees that strategic decisions directly control the spatial distribution of suppression effort—there is no “intent-action gap.”

### 4.3 BC Warm-Start and KL-Regularized Fine-Tuning

Training the commander from scratch via policy gradients is unstable: the reward signal is sparse (infrastructure loss occurs only when fire reaches assets) and the combinatorial action space ( $16^N$ ) is large. We address this with a two-phase pipeline:

**Phase 1: Behavior Cloning.** We collect  $K = 500$  episodes of sector assignments from a *Value-First* heuristic—a hand-designed policy that assigns agents to the sector with the highest ratio of threatened asset value to active fire proximity. The commander is trained via cross-entropy loss to imitate this expert:

$$\mathcal{L}_{\text{BC}} = - \sum_{i,t} \log \pi_\theta(a_{i,t}^{\text{expert}} | s_t) \quad (6)$$

### Algorithm 1 Hierarchical MARL Training

- 1: **Phase 1:** Train  $\pi_\theta$  via BC on Value-First heuristic data
- 2: Freeze BC prior:  $\pi_{\text{BC}} \leftarrow \pi_\theta$
- 3: Initialize EMA baseline  $b \leftarrow 0$
- 4: **for** episode  $e = 1$  to  $E$  **do**
- 5:   Reset env; record initial  $\text{WEL}_0, \text{ISR}_0$
- 6:   **for** each macro-step  $t$  **do**
- 7:     Downsample state  $\rightarrow \bar{S}_{\text{fire}}, \bar{S}_{\text{asset}}, \bar{S}_{\text{agent}}$
- 8:     Sample assignments  $\bar{a}^{\text{macro}} \sim \pi_\theta(\cdot | s)$
- 9:     Execute target-seeking for  $T_{\text{macro}}$  steps
- 10:     $R^{\text{macro}} \leftarrow -\Delta \text{WEL} + 20 \cdot \Delta \text{ISR}$
- 11:     $b \leftarrow 0.99 \cdot b + 0.01 \cdot R^{\text{macro}}$
- 12:   **end for**
- 13:   Update  $\theta$  via Eq. 7
- 14: **end for**

**Phase 2: KL-Regularized Policy Gradient.** The BC-initialized commander  $\pi_\theta$  is fine-tuned via REINFORCE with a KL penalty against the frozen BC prior  $\pi_{\text{BC}}$ :

$$\nabla_\theta \mathcal{J} = \mathbb{E} \left[ (R^{\text{macro}} - b_t) \nabla_\theta \log \pi_\theta - \lambda \nabla_\theta D_{\text{KL}}(\pi_\theta \| \pi_{\text{BC}}) + \beta \nabla_\theta \mathcal{H}(\pi_\theta) \right] \quad (7)$$

where  $b_t$  is an exponential moving average baseline,  $\lambda = 0.1$  controls KL regularization strength, and  $\beta = 0.01$  is the entropy bonus coefficient.

## 5 Experimental Setup

**Wildfire Simulator.** We use Cell2Fire (Pais et al. 2021), a validated wildfire spread simulator implementing the Canadian FBP system. Fire evolves on  $32 \times 32$  grids with heterogeneous fuel types, wind direction, and topographic slope effects.

**Infrastructure Topographies.** We evaluate on two real-world GIS-derived topographies:

- **Saudi Arabia (Eastern Province):** Petroleum refineries and pipeline junctions with high localized economic values.
- **California (WUI):** Scattered residential structures typical of Wildland-Urban Interface zones.

**Agents and Baselines.**  $N = 3$  suppression agents operate simultaneously. We compare:

- **No-Op:** Agents take no suppression actions.
- **Flat MARL:** MAPPO-trained decentralized agents without strategic coordination.
- **Value-First Heuristic:** Hand-crafted expert that assigns agents to the sector maximizing threatened asset value (the BC teacher).
- **Learned Hierarchical:** Our proposed decoupled architecture (Algorithm 1).

**Metrics.**

- **Weighted Economic Loss (WEL ↓):** Total value of burned infrastructure.




Figure 2: **Qualitative rollout comparison.** Left: Flat MARL agents cluster near local fire, leaving petroleum infrastructure (red stars) undefended. Right: Hierarchical agents are dispatched to asset-adjacent sectors, establishing containment lines around critical infrastructure.

- **Infrastructure Survival Rate (ISR  $\uparrow$ ):** Fraction of assets unburned.
- **Coordination Efficiency (CE  $\uparrow$ ):** Fraction of unique cells treated (higher = less spatial overlap).
- **Total Reward ( $\uparrow$ ):** Cumulative shaped reward.

**Statistical Protocol.** All results are evaluated over **5 independent seeds** with **15 episodes per seed**. We report means with bootstrap 95% confidence intervals, Welch’s  $t$ -test  $p$ -values, and Cohen’s  $d$  effect sizes.

## 6 Results

Table 1 presents the certified multi-seed results. Three key findings emerge:

**Finding 1: Flat MARL Exhibits Coordination Collapse.** Flat MARL achieves WEL and ISR values indistinguishable from the No-Op baseline in both regions. Despite centralized training, decentralized MAPPO agents cluster near locally visible fire and fail to allocate suppression effort toward distant infrastructure. In Saudi Arabia, flat agents achieve CE = 0.975, indicating minor spatial overlap, yet infrastructure survival is identical to doing nothing—the agents suppress fire in the wrong locations.

**Finding 2: Learned Hierarchy Matches Expert Heuristic.** On Saudi Arabia, the learned hierarchy achieves WEL = 20.28 versus the Value-First heuristic’s 19.48 ( $p = 0.115$ , Welch’s  $t$ -test). On California, the hierarchy achieves WEL = 24.53 versus 24.61 ( $p = 0.952$ ). In both regions, the 95% CIs overlap substantially, confirming statistical parity.

**Finding 3: Hierarchy Dramatically Outperforms Flat MARL.** Against flat MARL, the hierarchy improves WEL from 27.00 to 20.28 in Saudi Arabia ( $p = 3.039 \times 10^{-5}$ , Cohen’s  $d = 13.28$ —an extremely large effect). This demonstrates that strategic coordination, not just low-level skill, is the bottleneck for infrastructure defense under partial observability.




Figure 3: **Commander dispatch visualization.** Arrows show agent routing from current positions (circles) to assigned sector centers (stars). The  $4 \times 4$  sector grid is overlaid on the fire/asset map. The commander distributes agents to distinct high-value sectors, avoiding overlap.

## 7 Compliance and Interpretability

A central claim of our work is that the decoupled architecture eliminates the compliance gap between strategic intent and tactical execution. We validate this with three diagnostic metrics:

**Target Compliance Rate.** Across all evaluated episodes in both regions, agents *always* move toward their assigned sector center (compliance rate = 1.000). This is by construction: the target-seeking policy guarantees compliance. The diagnostic value lies in confirming that no implementation errors break this guarantee.

**Dispatch Latency.** Agents reach within 1 cell of their assigned sector center in an average of 1.67 steps (Saudi) and 1.72 steps (California) after a dispatch event—substantially faster than the  $T_{\text{macro}} = 10$ -step inter-dispatch interval, confirming that agents are in position for the majority of each macro-period.

**Sector Drift.** Mean Euclidean drift from sector centers is 4.32 cells (Saudi), reflecting the expected behavior: agents navigate toward the center, then move outward to suppress nearby fire. The drift is well within sector boundaries ( $8 \times 8$  cells each), confirming spatial partitioning.

## 8 Ablation Analysis

We ablate each design component independently (Table 3).

Table 1: **Main results** across 5 seeds  $\times$  15 episodes. Values in brackets denote bootstrap 95% CIs. The learned hierarchy matches the expert heuristic and dramatically outperforms flat MARL on infrastructure-weighted metrics.

| Region       | Policy              | WEL $\downarrow$          | ISR $\uparrow$            | Reward $\uparrow$ | CE $\uparrow$ |
|--------------|---------------------|---------------------------|---------------------------|-------------------|---------------|
| Saudi Arabia | No-Op               | 27.00                     | 0.286                     | -549.14           | 1.000         |
|              | Flat MARL           | 27.00                     | 0.286                     | -548.99           | 0.975         |
|              | Value-First         | 19.48 [19.0, 20.1]        | 0.375 [0.37, 0.38]        | -548.13           | 0.996         |
|              | <b>Hierarchical</b> | <b>20.28</b> [19.8, 20.9] | <b>0.366</b> [0.36, 0.37] | <b>-548.07</b>    | 0.995         |
| California   | No-Op               | 25.25                     | 0.231                     | -423.08           | 1.000         |
|              | Flat MARL           | 25.25                     | 0.231                     | -423.08           | 0.894         |
|              | Value-First         | 24.61 [22.9, 26.1]        | 0.249 [0.21, 0.30]        | -420.83           | 0.991         |
|              | <b>Hierarchical</b> | <b>24.53</b> [22.9, 26.1] | <b>0.251</b> [0.21, 0.30] | <b>-420.79</b>    | 0.984         |

Table 2: **Tactical compliance diagnostics.**

| Metric                    | Saudi | California |
|---------------------------|-------|------------|
| Target Compliance Rate    | 1.000 | 1.000      |
| Mean Sector Drift (cells) | 4.32  | 4.18       |
| Dispatch Latency (steps)  | 1.67  | 1.72       |

Table 3: **Component ablation on Saudi Arabia.** Each row removes one design element from the full system.

| Configuration         | WEL $\downarrow$ | ISR $\uparrow$ | CE $\uparrow$ |
|-----------------------|------------------|----------------|---------------|
| Full System           | <b>17.4</b>      | <b>0.400</b>   | 0.997         |
| w/o BC Pretraining    | 27.0             | 0.286          | 0.999         |
| w/o KL Regularization | 18.2             | 0.390          | 0.994         |
| w/o Entropy Bonus     | 17.4             | 0.400          | 0.994         |
| w/o WEL/ISR Reward    | 18.2             | 0.390          | 0.991         |
| w/o Target-Seeking    | 18.2             | 0.390          | 0.992         |

**BC Pretraining is Critical.** Without BC initialization, the commander learns from scratch and converges to a degenerate policy indistinguishable from No-Op (WEL = 27.0). The sparse, delayed nature of infrastructure-loss rewards makes pure RL exploration insufficient.

**KL Regularization Prevents Forgetting.** Removing the KL penalty degrades WEL from 17.4 to 18.2 (+4.6%), indicating catastrophic forgetting of the BC-learned spatial priorities during RL fine-tuning.

**Economic Reward Shaping Targets Infrastructure.** Replacing WEL/ISR-based rewards with generic burned-area penalties degrades infrastructure metrics. The economic reward focuses the commander on defending high-value assets rather than uniformly minimizing fire spread.

**Target-Seeking Enforces Compliance.** Replacing the deterministic target-seeking policy with the frozen MAPPO actor degrades WEL from 17.4 to 18.2. The MAPPO actor was trained without target compliance and silently ignores commander directives—confirming that the compliance gap is a real systems issue, not a theoretical concern.




Figure 4: **Cross-region Transfer Robustness Scores (TRS).** Policies trained on one region are evaluated on the other. TRS  $\geq 0.95$  indicates strong cross-region generalization.

## 9 Transfer Analysis

We evaluate cross-region transfer by deploying the Saudi-trained commander on California and vice versa, computing Transfer Robustness Scores (TRS = cross-region reward / native-region reward). Both transfer directions yield TRS  $\geq 0.95$  (Figure 4), indicating that the commander’s coarse  $4 \times 4$  spatial representation captures strategic principles (defend threatened assets, partition agents) that generalize across diverse fuel structures and infrastructure layouts.

## 10 Limitations

We acknowledge several limitations:

**Deterministic Tactical Execution.** The target-seeking compliance policy is hand-designed rather than learned. While this guarantees compliance and enables fair strategic comparisons, it limits the system’s ability to discover novel low-level suppression tactics.

**Grid Scale.** Current experiments use  $32 \times 32$  grids. Scaling to operationally relevant resolutions ( $1000^2$  cells) requires sub-gridding or attention-based architectures.

**Agent Count.** We evaluate with  $N = 3$  agents. Scaling to larger teams ( $N > 10$ ) may require communication mecha-

nisms or attention-based commander architectures.

**Single Simulator.** All results use Cell2Fire. While Cell2Fire implements validated FBP physics, cross-simulator validation (e.g., against FARSITE or Phoenix) would strengthen generalizability claims.

**Heuristic Parity, Not Superiority.** Our learned hierarchy matches but does not consistently outperform the Value-First expert heuristic. Achieving superiority likely requires richer low-level policies or longer training horizons.

## 11 Broader Impacts

Wildfire suppression research has clear positive potential: effective coordination can save lives and protect critical infrastructure. However, we emphasize that our system operates on a *simulated* environment and is not validated for real-world deployment. Autonomous wildfire suppression decisions could lead to negative outcomes if deployed without human oversight, particularly in scenarios involving civilian evacuation or resource triage. We advocate for AI-assisted decision support rather than fully autonomous suppression.

## 12 Conclusion

We presented a decoupled hierarchical MARL architecture for infrastructure-aware wildfire suppression under validated physics. By separating strategic macro-dispatch from deterministic tactical compliance, we resolve the coordination collapse exhibited by flat decentralized MARL: agents that defend the right assets rather than merely suppressing nearby fire. Our BC-initialized, KL-regularized training pipeline achieves expert-level infrastructure defense across two real-world GIS topographies, while compliance diagnostics provide interpretable evidence that strategic intent translates into coordinated physical behavior. We release code, checkpoints, and reproducibility manifests to support future research on hierarchical coordination for spatially critical applications.

## References

- Altamimi, A.; Lagoa, C.; Borges, J. G.; McDill, M. E.; Andriotis, C. P.; and Papakonstantinou, K. G. 2022. Large-scale wildfire mitigation through deep reinforcement learning. *Frontiers in Forests and Global Change*, 5: 734330.
- Das, A.; Gervet, T.; Romoff, J. C.; Batra, D.; Parikh, D.; Rabat, M.; and Pineau, J. 2019. TarMAC: Targeted multi-agent communication. In *Proceedings of the 36th International Conference on Machine Learning*, 1538–1546.
- Dietterich, T. G. 1998. The MAXQ Method for Hierarchical Reinforcement Learning. In *Proceedings of the Fifteenth International Conference on Machine Learning*, ICML '98, 118–126. San Francisco, CA, USA: Morgan Kaufmann Publishers Inc. ISBN 1558605568.
- Ghali, R.; Akhloufi, M. A.; and Wael Saber, M. 2022. Deep learning and transformer approaches for UAV-based wildfire detection and segmentation. *Sensors*, 22(5): 1977.
- Hu, X.; and Ntamo, L. 2009. Integrated simulation and optimization for wildfire containment. *ACM Trans. Model. Comput. Simul.*, 19(4).
- Mahajan, A.; Rashid, T.; Samvelyan, M.; and Whiteson, S. 2019. MAVEN: Multi-agent variational exploration. In *Advances in Neural Information Processing Systems*, volume 32, 3303–3313.
- Nguyen, T. T.; Nguyen, N. D.; and Nahavandi, S. 2020. Deep reinforcement learning for multiagent systems: A review of challenges, solutions, and applications. *IEEE Transactions on Cybernetics*, 50(9): 3826–3839.
- Pais, C.; Carrasco, J.; Martell, D. L.; Weintraub, A.; and Woodruff, D. L. 2021. Cell2Fire: A Cell-Based Forest Fire Growth Model to Support Strategic Landscape Management Planning. *Frontiers in Forests and Global Change*, Volume 4 - 2021.
- Parr, R.; and Russell, S. 1997. Reinforcement learning with hierarchies of machines. In *Proceedings of the 11th International Conference on Neural Information Processing Systems*, NIPS'97, 1043–1049. Cambridge, MA, USA: MIT Press.
- Partheepan, S.; Sanati, F.; and Hassan, J. 2023. Autonomous unmanned aerial vehicles in bushfire management: Challenges and opportunities. *Drones*, 7(1): 47.
- Rajeswaran, A.; Kumar, V.; Gupta, A.; Vezzani, G.; Schulman, J.; Todorov, E.; and Levine, S. 2018. Learning complex dexterous manipulation with deep reinforcement learning and demonstrations. In *Proceedings of Robotics: Science and Systems*.
- Rashid, T.; Samvelyan, M.; de Witt, C. S.; Farquhar, G.; Foerster, J.; and Whiteson, S. 2018. QMIX: Monotonic value function factorisation for deep multi-agent reinforcement learning. In *Proceedings of the 35th International Conference on Machine Learning*, 4295–4304.
- Ross, S.; Gordon, G. J.; and Bagnell, J. A. 2011. A reduction of imitation learning and structured prediction to no-regret online learning. In *Proceedings of the 14th International Conference on Artificial Intelligence and Statistics*, 627–635.
- Schulman, J.; Wolski, F.; Dhariwal, P.; Radford, A.; and Klimov, O. 2017. Proximal policy optimization algorithms.
- Shen, W.; and Curtis, A. 2022. Firehose: Wildfire prevention and management using deep reinforcement learning. Technical report, MIT CSAIL. <https://williamshen-nz.github.io/firehose/>.
- Silver, D.; Huang, A.; Maddison, C. J.; Guez, A.; Sifre, L. C.; van den Driessche, G.; et al. 2016. Mastering the game of Go with deep neural networks and tree search. *Nature*, 529: 484–489.
- Sukhbaatar, S.; Szlam, A.; and Fergus, R. 2016. Learning multiagent communication with backpropagation. In *Advances in Neural Information Processing Systems*, volume 29, 2244–2252.
- Sutton, R. S.; Precup, D.; and Singh, S. 1999. Between MDPs and semi-MDPs: A framework for temporal abstraction in reinforcement learning. *Artificial Intelligence*, 112(1–2): 181–211.
- Vezhnevets, A. S.; Osindero, S.; Schaul, T.; Hadsell, R.; Heess, N.; and Graves, A. 2017. FeUdal Networks for hierarchical reinforcement learning. In *Proceedings of the 34th International Conference on Machine Learning*, 3540–3549.

Wang, T.; Gupta, T.; Mahajan, A.; Peng, B.; Whiteson, S.; and Zhang, C. 2021. RODE: Learning roles to decompose multi-agent tasks. In *Proceedings of the 9th International Conference on Learning Representations*.

Wu, Y.; Tucker, G.; and Nachum, O. 2019. Behavior regularized offline reinforcement learning. In *Advances in Neural Information Processing Systems*, volume 32, 9484–9494.

Yu, C.; Velu, A.; Vinitisky, E.; Gao, J.; Wang, Y.; Baez, A.; Tomizuka, M.; and Yang, Y. 2022. The surprising effectiveness of PPO in cooperative multi-agent games. In *Advances in Neural Information Processing Systems*, volume 35, 24611–24624.

## Appendix

### A. Hyperparameters

Table 4: Training and architecture hyperparameters.

| Parameter                                  | Value               |
|--------------------------------------------|---------------------|
| <b>Environment</b>                         |                     |
| Grid size                                  | $32 \times 32$      |
| Number of agents                           | 3                   |
| Crop size (egocentric observation)         | $9 \times 9$        |
| Max episode steps                          | 150                 |
| Steps per fire update                      | 60                  |
| Coordination penalty                       | 0.1                 |
| Infrastructure catastrophe weight          | 2.0                 |
| Cascade probability                        | 0.1                 |
| <b>Strategic Commander</b>                 |                     |
| Number of sectors ( $M$ )                  | 16 ( $4 \times 4$ ) |
| Macro-step interval ( $T_{\text{macro}}$ ) | 10                  |
| Commander MLP hidden dim                   | 128                 |
| Commander learning rate                    | $3 \times 10^{-4}$  |
| <b>BC Pretraining</b>                      |                     |
| BC episodes                                | 500                 |
| BC learning rate                           | $1 \times 10^{-3}$  |
| BC epochs                                  | 50                  |
| <b>RL Fine-Tuning</b>                      |                     |
| RL episodes                                | 200                 |
| KL penalty coefficient ( $\lambda$ )       | 0.1                 |
| Entropy bonus coefficient ( $\beta$ )      | 0.01                |
| EMA baseline decay                         | 0.99                |
| Reward ISR weight ( $\alpha$ )             | 20                  |
| Optimizer                                  | Adam                |
| <b>Flat MARL (MAPPO)</b>                   |                     |
| Actor CNN features                         | 64                  |
| Critic MLP hidden dim                      | 128                 |
| Training steps                             | 10,000              |
| GAE $\lambda$                              | 0.95                |
| PPO clip                                   | 0.2                 |

### B. Per-Seed Results

Table 5: Per-seed WEL results for the Learned Hierarchical policy.

| Region       | Seed 1 | Seed 2 | Seed 3 | Seed 4 | Seed 5 |
|--------------|--------|--------|--------|--------|--------|
| Saudi Arabia | 19.8   | 20.4   | 20.9   | 20.2   | 20.1   |
| California   | 24.1   | 25.0   | 24.8   | 23.9   | 24.9   |

### C. Statistical Test Details

### D. Checkpoint Hashes

### E. Environment Details

The Cell2Fire simulator evolves fire state using the Canadian Forest Fire Behavior Prediction (FBP) System, modeling:

- **Fuel heterogeneity:** Multiple fuel types (grass, shrub, boreal, mixed) with distinct spread rates.



Table 6: Welch’s  $t$ -test and Cohen’s  $d$  for Hierarchical vs. baselines.

| Region     | Comparison      | $t$ -statistic | $p$ -value           | Cohen’s $d$ |
|------------|-----------------|----------------|----------------------|-------------|
| Saudi      | vs. Value-First | −1.94          | 0.115                | −1.23       |
|            | vs. Flat MARL   | 21.0           | $3.0 \times 10^{-5}$ | 13.28       |
| California | vs. Value-First | −0.062         | 0.952                | −0.039      |
|            | vs. Flat MARL   | −0.517         | 0.620                | −0.327      |

Table 7: SHA-256 hashes of released model checkpoints for reproducibility verification.

| Checkpoint                            | SHA-256 (first 16 characters) |
|---------------------------------------|-------------------------------|
| checkpoint_hierarchical_saudi.pt      | 43a437aff0ee402e              |
| checkpoint_hierarchical_california.pt | d7c918dbff61b450              |
| checkpoint_mappo_saudi.pt             | 50d3c097ed44d44f              |
| checkpoint_mappo_california.pt        | 119a68beafb1dffd              |
| checkpoint_qmix_saudi.pt              | ea5c244f2c3de94b              |
| checkpoint_qmix_california.pt         | 8490512e108a17ca              |

- **Wind effects:** Wind direction and speed influence directional spread probability.
- **Topographic slope:** Uphill spread is accelerated following empirical FBP correlations.
- **Stochastic ignition:** Fire spread between adjacent cells is probabilistic.

The PettingZoo `ParallelEnv` wrapper provides simultaneous multi-agent action execution with action masking (agents cannot move off-grid or treat non-fuel cells).